

DSBDA-5 - Jupyter Notebook

localhost:8889/notebooks/DSBDA/DSBDA-5.ipynb

jupyter DSBDA-5 Last Checkpoint: 7 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [4]: import matplotlib.pyplot as plt
import pandas as pd

df = pd.read_csv('Social_Network_Ads.csv')
df.head()
```

Out[4]:

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0

```
In [5]: x = df[['Age', 'EstimatedSalary']]
y = df['Purchased']

In [6]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
X_train.shape
```

Out[6]: (300, 2)

```
In [7]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
```

Out[7]: LogisticRegression()

```
In [8]: model.score(X_test, y_test)
```

DSBDA-5 - Jupyter Notebook

localhost:8889/notebooks/DSBDA/DSBDA-5.ipynb

jupyter DSBDA-5 Last Checkpoint: 9 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [8]: model.score(X_test, y_test)
```

Out[8]: 0.63

```
In [9]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [10]: model = LogisticRegression()
model.fit(X_train_scaled, y_train)
```

Out[10]: LogisticRegression()

```
In [11]: model.score(X_test_scaled, y_test)
```

Out[11]: 0.84

```
In [12]: X_train_scaled
```

Out[12]: array([[0.5, 0.6],
[0.76190476, 0.15555556],
[0.5, 0.41481481],
[0.0952381, 0.35555556],
[0.4047619, 0.32592593],
[0.92857143, 0.08148148],
[0.19047619, 0.11111111],
[0.4047619, 0.17777778],
[0.61904762, 0.91851852],
[0.45238095, 0.97037037],
[0.14285714, 0.12592593],
[0.21428571, 0.01481481],
[0.45238095, 0.57777778],
[0.4047619, 0.07407407],
[0.19047619, 0.48888889],

DSBDA-5 - Jupyter Notebook

localhost:8889/notebooks/DSBDA/DSBDA-5.ipynb

jupyter DSBDA-5 Last Checkpoint: 9 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [12]: `x_train_scaled`

```
Out[12]: array([[0.5, 0.6],
 [0.76190476, 0.15555556],
 [0.5, 0.41481481],
 [0.0952381, 0.35555556],
 [0.4047619, 0.32592593],
 [0.92857143, 0.08148148],
 [0.19047619, 0.11111111],
 [0.4047619, 0.17777778],
 [0.61904762, 0.91851852],
 [0.45238095, 0.97037037],
 [0.14285714, 0.12592593],
 [0.21428571, 0.01481481],
 [0.45238095, 0.57777778],
 [0.4047619, 0.07407407],
 [0.19047619, 0.48888889],
 [0., 0.4962963 ],
 [0.33333333, 0.02222222],
 [0.42857143, 0.25925926],
 [0.30952381, 0.43703704],
 [0.30952381, 0.32592593]])
```

In [13]: `from sklearn.metrics import confusion_matrix`
`y_p = model.predict(X_test_scaled)`
`y_p`

```
Out[13]: array([0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1,
 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1,
 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0,
 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0], dtype=int64)
```

In [14]: `confusion_matrix(y_test, y_p)`

```
Out[14]: array([[58, 5],
 [11, 26]], dtype=int64)
```